Not Just A Black Box:
Interpretable Deep Learning by Propagating Activation Differences

Avanti Shrikumar[1] (avanti@stanford.edu), **Peyton Greenside**[2] (pgreens@stanford.edu)
Anna Y. Shcherbina[2] (annashch@stanford.edu), **Anshul Kundaje**[1,3] (akundaje@stanford.edu)

1. Department of Computer Science, Stanford University, CA, USA
2. Biomedical Informatics, Stanford University, CA, USA
3. Department of Genetics, Stanford University, CA, USA

Abstract

The purported "black box" nature of neural networks is a barrier to adoption in applications where interpretability is essential. Here we present DeepLIFT (Learning Important Features), an efficient and effective method for computing importance scores in a neural network. DeepLIFT compares the activation of each neuron to its 'reference activation' and assigns contribution scores according to the difference. We apply DeepLIFT to models trained on natural images and genomic data, and show significant advantages over gradient-based methods.

1. Introduction

As neural networks become increasingly popular, their "black box" reputation is a barrier to adoption when interpretability is paramount. Understanding the features that lead to a particular output builds trust with users and can lead to novel scientific discoveries. Simonyan et al. (2013) proposed using gradients to generate saliency maps and showed that this is a generalization of the deconvolutional nets of Zeiler et al. (2013). Guided backpropagation (Springenberg et al. 2014) is another variant which only considers gradients that have positive error signal. As shown in Figure 2, saliency maps can be substantially improved by simply multiplying the gradient with the input signal, which corresponds to a first-order Taylor approximation of how the output would change if the input were set to zero; the layerwise relevance propagation rules described in Samek et al. (2015) reduce to this approach, assuming bias terms are included in the denominators.

Gradient-based approaches are problematic because activation functions such as Rectified Linear Units (ReLUs) have a gradient of zero when they are not firing, and yet a ReLU that does not fire can still carry information (Figure 1). Similarly, sigmoid or tanh activations are popular choices for the activation functions of gates in memory units of recurrent neural networks such as GRUs and LSTMs (Chung et al. 2014; Hochreiter and Schmidhuber 1997), but these activations have a near-zero gradient at high or low inputs even though such inputs can be very significant.
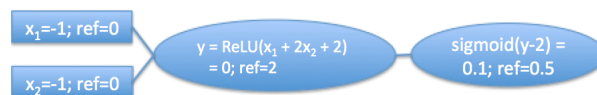

Figure 1. Simple network with inputs x_1 and x_2 that have reference values of 0. When $x_1 = x_2 = -1$, output is 0.1 but the gradients w.r.t x_1 and x_2 are 0 due to inactive ReLU y (which has activation of 2 under reference input). By comparing activations to their reference values, DeepLIFT assigns contributions to the output of $\left((0.1 - 0.5)\frac{1}{3}\right)$ to x_1 and $\left((0.1 - 0.5)\frac{2}{3}\right)$ to x_2.

We present DeepLIFT, a method for assigning feature importance that compares a neuron's activation to its 'reference', where the reference is the activation that the neuron has when the network is provided a 'reference input' (the reference input is defined according to what is appropriate for the task at hand). This addresses the limitation of gradient-based approaches because the difference from the reference may be non-zero even when the gradient is zero.

2. DeepLIFT Method

We denote the contribution of x to y as C_{xy}. Let the activation of a neuron n be denoted as A_n. Further, let the *reference* activation of neuron n be denoted A_n^0, and let the $A_n - A_n^0$ be denoted as δ_n. We define our contributions C_{xy} to satisfy the following properties.

2.1. Summation to δ

For any set of neurons S whose activations are minimally sufficient to compute the activation of y (that is, if we know the activations of S, we can compute the activation of y, and there is no set $S' \subset S$ such that S' is sufficient to compute the activation of y - in layman's terms, S is a full set of non-redundant inputs to y), the following property holds:

$$\sum_{s \in S} C_{sy} = \delta_y \tag{1}$$

That is, the sum over all the contributions of neurons in S to y equals the difference-from-reference of y.

2.2. Linear composition

Let O_x represent the output neurons of x. The following property holds:

$$C_{xy} = \sum_{o \in O_x} \frac{C_{xo}}{\delta_o} C_{oy} \tag{2}$$

In layman's terms, each neuron 'inherits' a contribution through its outputs in proportion to how much that neuron contributes to the difference-from-reference of the output.

2.3. Backpropagation Rules

We show that the contributions as defined above can be computed using the following rules (which can be implemented to run on a GPU). The computation is reminiscent of the chain rule used during gradient backpropagation, as equation 2 makes it possible to start with contribution scores of later layers and use them to find the contribution scores of preceding layers. To avoid issues of numerical stability when δ_n for a particular neuron is small, rather than computing the contribution scores explicitly, we instead compute *multipliers* m_{xy} that, when multiplied with the difference-from-reference, give the contribution:

$$m_{xy}\delta_x = C_{xy} \tag{3}$$

Let t represent the target neuron that we intend to compute contributions to, and let O_x represent the set of outputs of x. We show that:

$$m_{xt} = \sum_{y \in O_x} m_{xy} m_{yt} \tag{4}$$

The equation above follows from the linear composition property and the definition of the multipliers, as proved below:

$$C_{xt} = \sum_{y \in O_x} \frac{C_{xy}}{\delta_y} C_{yt}$$

$$m_{xt}\delta_x = \sum_{y \in O_x} \frac{C_{xy}}{\delta_y}(m_{yt}\delta_y) = \sum_{y \in O_x} C_{xy} m_{yt} \tag{5}$$

$$m_{xt} = \sum_{y \in O_x} \frac{C_{xy}}{\delta_x} m_{yt} = \sum_{y \in O_x} m_{xy} m_{yt}$$

In the equations below, I_y denotes the set of inputs of y.

2.3.1. AFFINE FUNCTIONS

Let

$$A_y = \left(\sum_{x \in I_y} w_{xy} A_x\right) + b \tag{6}$$

Then $m_{xy} = w_{xy}$

Proof. We show that $\delta_y = \sum_{x \in I_y} m_{xy}\delta_x$.

Using the fact that $A_n = A_n^0 + \delta_n$, we have:

$$
\begin{aligned}
(A_y^0 + \delta_y) &= \left(\sum_{x \in I_y} w_{xy}(A_x^0 + \delta_x)\right) + b \\
&= \left(\sum_{x \in I_y} w_{xy} A_x^0\right) + b + \sum_{x \in I_y} w_{xy}\delta_x
\end{aligned}
\tag{7}
$$

We also note that the reference activation A_y^0 can be found as follows:

$$A_y^0 = \left(\sum_{x \in I_y} w_{xy} A_x^0\right) + b \tag{8}$$

Thus, canceling out A_y^0 yields:

$$\delta_y = \sum_{x \in I_y} w_{xy}\delta_x = \sum_{x \in I_y} m_{xy}\delta_x \tag{9}$$

2.3.2. MAX OPERATION

We consider the case of max operation such as a maxpool:

$$A_y = \max_{x \in I_y} A_x \tag{10}$$

Then we have:

$$m_{xy} = \mathbf{1}\{A_x = A_y\}\frac{\delta_y}{\delta_x} \tag{11}$$

Where $\mathbf{1}\{\}$ is the indicator function. If a symbolic computation package is used, then the gradient of y with respect to x can be used in place of $\mathbf{1}\{A_x = A_y\}$.

Proof.

$$\sum_{x \in y} m_{xy} \delta_x = \left(\sum_{x \in y} \mathbf{1}\{A_x = A_y\} \frac{\delta_y}{\delta_x} \right) \delta_x$$
$$= \sum_{x \in y} \mathbf{1}\{A_x = A_y\} \delta_y = \delta_y \tag{12}$$

2.3.3. MAXOUT UNITS

A maxout function has the form

$$A_y = \max_{i=1}^{n} \left(\sum_x w_{xy}^i A_x \right) + b^i \tag{13}$$

i.e. it is the max over n affine functions of the input vector $\vec{x}$. For a given vector of activations $A_{\vec{x}}$ of the inputs, we split $A_{\vec{x}} - A_{\vec{x}}^0$ into segments such that over each segment s, a unique affine function dominates the maxout and the coefficient of an individual input x over that segment is $w(s)_{xy}$. Let $l(s)$ denote the fraction of $A_{\vec{x}} - A_{\vec{x}}^0$ in segment s. We have:

$$m_{xy} = \sum_s l(s) w(s)_{xy} \tag{14}$$

Intuitively speaking, we simply split the piecewise-linear maxout function into regions where it is linear, and do a weighted sum of the coefficients of x in each region according to how much of $A_{\vec{x}} - A_{\vec{x}}$ falls in that region.

2.3.4. OTHER ACTIVATIONS

The following choice for m_{xy}, which is the same for all inputs to y, satisfies summation-to-delta:

$$m_{xy} = \frac{\delta_y}{\sum_{x' \in I_y} \delta_{x'}} \tag{15}$$

This rule may be used for nonlinearities like ReLUs, PReLUs, sigmoid and tanh (where y has only one input). Situations where the denominator is near zero can be handled by applying L'hopital's rule, because by definition:

$$\delta_y \to 0 \text{ as } \sum_{x \in I_y} \delta_x \to 0 \tag{16}$$

2.3.5. ELEMENT-WISE PRODUCTS

Consider the function:

$$A_y = A_y^0 + \delta_y = (A_{x_1}^0 + \delta_{x_1})(A_{x_2}^0 + \delta_{x_2}) \tag{17}$$

We have:

$$\delta_y = (A_{x_1}^0 + \delta_{x_1})(A_{x_2}^0 + \delta_{x_2}) - (A_{x_1}^0 A_{x_2}^0)$$
$$= A_{x_1}^0 \delta_{x_2} + A_{x_2}^0 \delta_{x_1} + \delta_{x_1} \delta_{x_2}$$
$$= \delta_{x_1} \left(A_{x_2}^0 + \frac{\delta_{x_2}}{2} \right) + \delta_{x_2} \left(A_{x_1}^0 + \frac{\delta_{x_1}}{2} \right) \tag{18}$$

Thus, viable choices for the multipliers are $m_{x_1 y} = A_{x_2}^0 + 0.5\delta_{x_2}$ and $m_{x_2 y} = A_{x_1}^0 + 0.5\delta_{x_1}$

2.4. A note on final activation layers

Activation functions such as a softmax or a sigmoid have a maximum δ of 1.0. Due to the *summation to δ* property, the contribution scores for individual features are lower when there are several redundant features present. As an example, consider $A_t = \sigma(A_y)$ (where $sigma$ is the sigmoid transformation) and $A_y = A_{x_1} + A_{x_2}$. Let the default activations of the inputs be $A_{x_1}^0 = A_{x_2}^0 = 0$. When $x_1 = 100$ and $x_2 = 0$, we have $C_{x_1 t} = 0.5$. However, when both $x_1 = 100$ and $x_2 = 100$, we have $C_{x_1 t} = C_{x_2 t} = 0.25$. To avoid this attenuation of contribution in the presence of redundant inputs, we can use the contributions to y rather than t; in both cases, $C_{x_1 y} = 100$.

2.5. A note on Softmax activation

Let $t_1, t_2 ... t_n$ represent the output of a softmax transformation on the nodes $y_1, y_2 ... y_n$, such that:

$$A_{t_i} = \frac{e^{A_{y_i}}}{\sum_{i'=1}^n e^{A_{y_i'}}} \tag{19}$$

Here, $A_{y_1} ... A_{y_n}$ are affine functions of their inputs. Let x represent a neuron that is an input to $A_{y_1} ... A_{y_n}$, and let w_{xy_i} represent the coefficient of A_x in A_{y_i}. Because $A_{y_1} ... A_{y_n}$ are followed by a softmax transformation, if w_{xy_i} is the same for all y_i (that is, x contributes equally to all y_i), then x effectively has zero contribution to A_{t_i}. This can be observed by substituting $A_{y_i} = w_{xy_i} A_x + r_{y_i}$ in the expression for A_{t_i} and canceling out $e^{w_{xy_i} A_x}$ (here, r_{y_i} is the sum of all the remaining terms in the affine expression for A_{y_i})

$$A_{t_i} = \frac{e^{A_{y_i}}}{\sum_{i'=1}^n e^{A_{y_i'}}} = \frac{e^{w_{xy_i} A_x + r_{y_i}}}{\sum_{i'=1}^n e^{w_{xy_{i'}} A_x + r_{y_{i'}}}}$$
$$= \frac{e^{w_{xy_i} A_x + r_{y_i}}}{\sum_{i'=1}^n e^{w_{xy_i} A_x + r_{y_{i'}}}} = \frac{e^{r_{y_i}}}{\sum_{i'=1}^n e^{r_{y_{i'}}}} \tag{20}$$

As mentioned in the previous subsection, in order to avoid attenuation of signal for highly confident predictions, we should compute C_{xy_i} rather than C_{xt_i}. One way to ensure that C_{xy_i} is zero if w_{xy_i} is the same for all y_i is to mean-normalized the weights as follows:

$$\bar{w}_{xy_i} = w_{xy_i} - \frac{1}{n} \sum_{i'=1}^n w_{xy_{i'}} \tag{21}$$

This transformation will not affect the output of the softmax, but will ensure that the DeepLIFT scores are zero when a particular node contributes equally to all softmax classes.

2.6. Weight normalization for constrained inputs

Let y be a neuron with some subset of inputs S_y that are constrained such that $\sum_{x \in S_y} A_x = c$ (for example, one-hot

encoded input satisfies the constraint $\sum_{x \in S_y} A_x = 1$, and a convolutional neuron operating on one-hot encoded rows has one constraint per column that it sees). Let the weights from x to y be denoted w_{xy} and let b_y be the bias of y. It is advisable to use normalized weights $\bar{w}_{xy} = w_{xy} - \mu$ and bias $\bar{b}_y = b_y + c\mu$, where μ is the mean over all w_{xy}. We note that this maintains the output of the neural net because, for any constant μ:

$$
\begin{aligned}
A_y &= \left(\sum A_x(\bar{w}_{xy} - \mu)\right) + (b_y + c\mu) \\
&= \left(\sum A_x w_{xy}\right) - \left(\sum A_x \mu\right) + (b_y + c\mu) \\
&= \left(\sum A_x w_{xy}\right) - c\mu + (b_y + c\mu) \\
&= \left(\sum A_x w_{xy}\right) + b_y
\end{aligned}
\tag{22}
$$

The normalization is desirable because, for affine functions, the multipliers m_{xy} are equal to the weights w_{xy} and are thus sensitive to μ. To take the example of a convolutional neuron operating on one-hot encoded rows: by mean-normalizing w_{xy} for each column in the filter, one can ensure that the contributions C_{xy} from some columns are not systematically overestimated or underestimated relative to the contributions from other columns.

3. Results

3.1. Tiny ImageNet

A model with the VGG16 (Long et al., 2015) architecture was trained using the Keras framework (Chollet, 2015) on a scaled-down version of the Imagenet dataset, dubbed 'Tiny Imagenet'. The images were 64×64 in dimension and belonged to one of 200 output classes. Results shown in Figure 2; the reference input was an input of all zeros after preprocessing.

3.2. Genomics

We apply DeepLIFT to models trained on genomic sequence. The positive class requires that the DNA patterns 'GATA' and 'CAGATG' appear in the length-200 sequence together. The negative class has only one of the two patterns appearing once or twice. Outside the core patterns (which were sampled from a generative model) we randomly sample the four bases A, C, G and T. A CNN was trained using the Keras framework (Chollet, 2015) on one-hot encoded sequences with 20 convolutional filters of length 15 and stride 1 and a max pool layer of width and stride 50, followed by two fully connected layers of size 200. PReLU nonlinearities were used for the hidden layers. This model performs well with auROC of 0.907. The misclassified examples primarily occur when one of the patterns erroneously arises in the randomly sampled background.


Figure 2. Comparison of methods. Importance scores for RGB channels were summed to get per-pixel importance. Left-to-right: original image, absolute value of the gradient (similar to Simonyan et al. which used the two-norm across RGB rather than sum, and which generalizes Zeiler et al.), positive gradient*input (Taylor approximation, equivalent to Layerwise Relevance Propagation in Samek et al. but masking negative contributions as described in Springenberg et al.), and positive DeepLIFT.

We then run DeepLIFT to assign an importance score to each base in the correctly predicted sequences (the reference input for DeepLIFT was an input of all zeros post weight-normalization (see 2.6) of the first conv layer) and compared the results to the gradient*input (Figure 3).


Figure 3. DeepLIFT scores (top) and gradient*input (bottom) are plotted for each position in the DNA sequence and colored by the DNA base (due to one-hot encoding, input is either 1 or 0; gradient*input is equivalent to taking the gradient for the letter that is actually present). DeepLIFT discovers both patterns and assigns them large importance scores. Gradient-based methods miss the GATA pattern.

3.3. Discussion

Prevailing feature importance methods such as the saliency maps of Simonyan et al. (which generalize the deconvolutional nets of Zeiler et al.) and the guided backpropagation of Springenberg et al. are variants of computing gradients. As shown in Figure 1, this can give misleading results when the local gradient is zero. DeepLIFT instead considers the deviation from a neuron's reference activity. This makes it capable of handling RNN memory units gated by activations that have vanishing gradients (eg: sigmoid, tanh).

Layerwise Relevance Propagation (LRP), described in Samek et al. and first proposed by Bach et al. (2015), does not obviously rely on gradients; however, it can be shown that if bias terms are included in the relevance propaga-

tion and all activations are piecewise linear, LRP reduces to gradient*input (a first-order Taylor approximation of the change in output if the input is set to zero). If all reference activations are zero (as happens when all bias terms are zero and all reference input values are zero), DeepLIFT and LRP give similar results (except that by computing contributions using multipliers, DeepLIFT circumvents the numerical stability problems that LRP faces). In practice, biases are often non-zero, which is why DeepLIFT produces superior results (Figures 2 & 3).

4. Author contributions

AS & PG conceived of DeepLIFT. AS implemented DeepLIFT in software. PG led application to genomics. AYS led application to Tiny Imagenet. AK provided guidance and feedback. AS, PG, AYS & AK prepared the manuscript.